

mini project

In [1]:

```
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.layers import *
from tensorflow import keras
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.utils import plot_model
import tensorflow as tf

from sklearn.model_selection import train_test_split

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

In [2]:

```
def plot_loss_acc(model):
    his_df = pd.DataFrame(model.history)

    fig, axes = plt.subplots(1, 2, figsize=(12, 4))
    axes[0].plot(his_df[['acc', 'val_acc']], label = ['Training', 'Validataion'])

    axes[0].set_title('Accarcy')
    axes[0].grid(True)
    axes[0].set_xlabel('epoch')

    axes[1].plot(his_df[['loss', 'val_loss']], label = ['Training',
'Validataion'])
    axes[1].set_title('loss')
    axes[1].grid(True)
    axes[1].set_xlabel('epoch')
    plt.show()
```

In [3]:

```
tr_dir = './data/Fast Food Classification V2/Train'
val_dir = './data/Fast Food Classification V2/Valid'
te_dir = './data/Fast Food Classification V2/Test'

image_gen = ImageDataGenerator(rescale = (1/255.))
tr_gen = image_gen.flow_from_directory(tr_dir,
                                      batch_size = 32,
                                      target_size = (244, 244),
                                      class_mode = 'categorical',
                                      seed = 2020)

val_gen = image_gen.flow_from_directory(val_dir,
                                       batch_size = 32,
                                       target_size = (244, 244),
                                       class_mode = 'categorical',
                                       seed = 2020)

te_gen = image_gen.flow_from_directory(te_dir,
                                       batch_size = 32,
                                       target_size = (244, 244),
                                       class_mode = 'categorical',
                                       seed = 2020)
```

Found 15000 images belonging to 10 classes.

Found 3500 images belonging to 10 classes.

Found 1500 images belonging to 10 classes.

In [4]:

```
class_labels= ['Baked Potato', 'Burger', 'Crispy Chicken', 'Donut', 'Fries', 'Hot  
Dog', 'Pizza', 'Sandwich', 'Taco', 'Taquito']
```

In [7]:

```
batch = next(tr_gen)
images, labels = batch[0], batch[1]
```

In [8]:

```
class_labels= ['Baked Potato', 'Burger', 'Crispy Chicken', 'Donut', 'Fries', 'Hot  
Dog', 'Pizza', 'Sandwich', 'Taco', 'Taquito']  
batch = next(tr_gen)  
images, labels = batch[0], batch[1]  
  
plt.figure(figsize=(15, 10))  
for i in range(24):  
    plt.subplot(4, 6, i + 1)  
    plt.imshow(images[i])  
    plt.axis('off')  
    plt.title(class_labels[np.argmax(labels, axis = 1)[i]])  
plt.subplots_adjust(wspace=0.5, hspace=0.5)  
plt.show()
```

Crispy Chicken



Burger



Fries




Burger



Pizza



Crispy Chicken



Taquito



Baked Potato



Pizza



Crispy Chicken



Fries



Donut



Taco



Sandwich



Burger



Hot Dog



Burger



Crispy Chicken




Fries



Taco



Crispy Chicken



Crispy Chicken



Burger



Pizza



In [9]:

```
from tensorflow.keras.applications import VGG16

pre_trained_base = VGG16(include_top = False,
                           weights="imagenet",
                           input_shape = [224, 224, 3])


pre_trained_base.trainable = False

plot_model(pre_trained_base, show_shapes = True, show_layer_names= True, to_file
= 'vgg16.png')
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5 (https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5)

58889256/58889256 [=====] - 3s 0us/step

Out[9]:



In [20]:

```
class vgg_bulid(Model):
    def __init__(self, pre_trained_base):
        super(vgg_bulid, self).__init__()
        self.pre_vgg16 = pre_trained_base
        self.flatten = tf.keras.layers.Flatten()
        self.fc1 = Sequential([
            tf.keras.layers.Dense(256,
                                   kernel_initializer = 'he_uniform'),
            tf.keras.layers.LeakyReLU(alpha = 0.1),
            tf.keras.layers.Dropout(0.5)
        ])
        self.fc2 = Sequential([
            tf.keras.layers.Dense(256,
                                   kernel_initializer = 'he_uniform'),
            tf.keras.layers.LeakyReLU(alpha = 0.2),
            tf.keras.layers.Dropout(0.5)
        ])
        self.soft_max = tf.keras.layers.Dense(10, activation = 'softmax')

    def call(self, x):
        ret = self.pre_vgg16(x)
        ret = self.flatten(ret)
        ret = self.fc1(ret)
        ret = self.fc2(ret)
        return self.soft_max(ret)
```

In [21]:

```
vgg_model = vgg_bulid(pre_trained_base)
vgg_model(tf.keras.Input(shape = (224, 224, 3)))
```

Out[21]:

<KerasTensor: shape=(None, 10) dtype=float32 (created by layer 'vgg_bulid_1')>

In [22]:

```
vgg_model.compile(optimizer = tf.keras.optimizers.Adam(learning_rate=0.001,
beta_1=0.9, beta_2=0.999,)),
                  loss = 'categorical_crossentropy',
                  metrics = ['acc'])
```

In [23]:

```
import datetime as dt

earlystopping = tf.keras.callbacks.EarlyStopping(monitor = 'val_loss', patience = 3)

checkpoint = tf.keras.callbacks.ModelCheckpoint(filepath =
f'/content/drive/MyDrive/4조/Team
Project/sojung/vgg_{dt.datetime.now().ctime()}.ckp',
                                                save_weights_only = True,
                                                save_best_only = True,
                                                monitor = 'val_loss',
                                                verbose = 1)
```


In [24]:

```
vgg_his = vgg_model.fit(tr_gen,
                        validation_data = val_gen,
                        epochs = 30,
                        callbacks = [checkpoint])
```

```
469/469 [=====] - 97s 206ms/step - loss: 0.3927 - acc: 0.8837
cc: 0.8665 - val_loss: 1.3452 - val_acc: 0.6583
Epoch 18/30
469/469 [=====] - ETA: 0s - loss: 0.4233 - acc: 0.8715
Epoch 18: val_loss did not improve from 1.10976
469/469 [=====] - 100s 213ms/step - loss: 0.4233 - acc: 0.8715 - val_loss: 1.4876 - val_acc: 0.6820
Epoch 19/30
469/469 [=====] - ETA: 0s - loss: 0.4063 - acc: 0.8782
Epoch 19: val_loss did not improve from 1.10976
469/469 [=====] - 97s 206ms/step - loss: 0.4063 - acc: 0.8782 - val_loss: 1.3988 - val_acc: 0.6660
Epoch 20/30
469/469 [=====] - ETA: 0s - loss: 0.3927 - acc: 0.8837
Epoch 20: val_loss did not improve from 1.10976
469/469 [=====] - 97s 206ms/step - loss: 0.3927 - acc: 0.8837 - val_loss: 1.6391 - val_acc: 0.6563
Epoch 21/30
```

In [28]:

```
plot_loss_acc(vgg_his)
```



In [29]:

```
vgg_model.evaluate(te_gen)
```

47/47 [=====] - 11s 244ms/step - loss: 1.5956 - acc: 0.7013

Out[29]:

```
[1.5955734252929688, 0.7013333439826965]
```

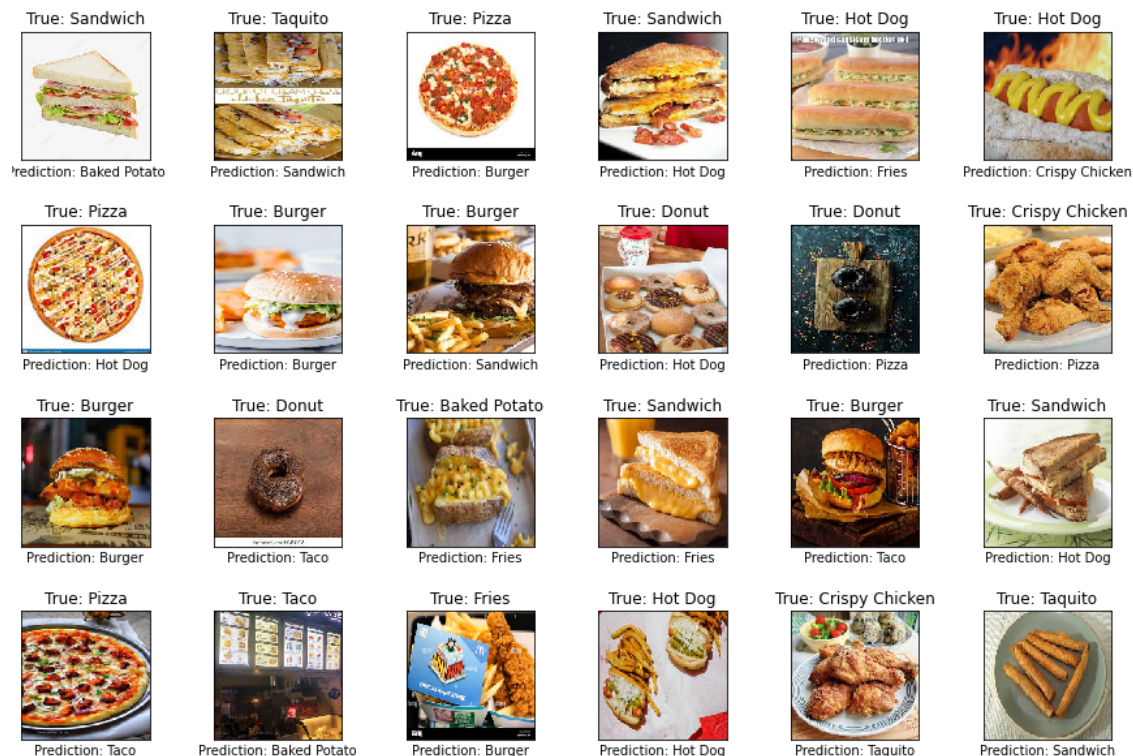
In [30]:

```
y_pred = np.argmax(vgg_model.predict(te_gen), axis = -1)

batch = next(tr_gen)
images, labels = batch[0], batch[1]

plt.figure(figsize=(15, 10))
for i in range(24):
    plt.subplot(4, 6, i + 1)
    plt.imshow(images[i], interpolation = 'nearest')
    # plt.axis('off')
    plt.xticks([])
    plt.yticks([])
    plt.xlabel(f'Prediction: {class_labels[y_pred[i]]}')
    plt.title('True: ' + class_labels[np.argmax(labels, axis = 1)[i]])
plt.subplots_adjust(wspace=0.5, hspace=0.5)
plt.show()
```

47/47 [=====] - 9s 184ms/step



In []:

In []:

가중치를 imagenet이 아닌 내가 학습시킴

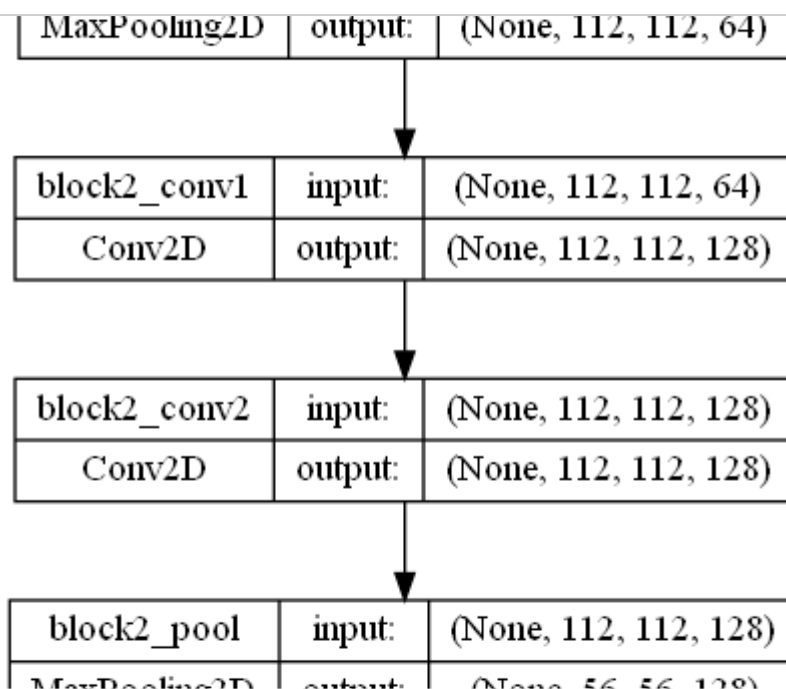
In [14]:

```
from tensorflow.keras.applications import VGG16

pre_trained_base = VGG16(include_top = False,
                          weights="imagenet",
                          input_shape = [224, 224, 3])

pre_trained_base.trainable = False

plot_model(pre_trained_base, show_shapes = True, show_layer_names= True, to_file
= 'vgg16.png')
```



In [15]:

```
class vgg_bulid(Model):
    def __init__(self, pre_trained_base):
        super(vgg_bulid, self).__init__()
        self.pre_vgg16 = pre_trained_base
        self.flatten = tf.keras.layers.Flatten()
        self.fc1 = Sequential([
            tf.keras.layers.Dense(256,
                                   kernel_initializer = 'he_uniform'),
            tf.keras.layers.LeakyReLU(alpha = 0.1),
            tf.keras.layers.BatchNormalization(),
            tf.keras.layers.Dropout(0.5)
        ])
        self.fc2 = Sequential([
            tf.keras.layers.Dense(256,
                                   kernel_initializer = 'he_uniform'),

            tf.keras.layers.LeakyReLU(alpha = 0.2),
            tf.keras.layers.BatchNormalization(),
            tf.keras.layers.Dropout(0.5)
        ])
        self.soft_max = tf.keras.layers.Dense(10, activation = 'softmax')

    def call(self, x):
        ret = self.pre_vgg16(x)
        ret = self.flatten(ret)
        ret = self.fc1(ret)
        ret = self.fc2(ret)
        return self.soft_max(ret)
```

In [16]:

```
vgg_model2 = vgg_bulid(vgg16_not_weight)
vgg_model2(tf.keras.Input(shape = (224, 224, 3)))
vgg_model2.summary()
```

Model: "vgg_bulid_2"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 7, 7, 512)	14714688
flatten_2 (Flatten)	multiple	0
sequential_4 (Sequential)	(None, 256)	6423808
sequential_5 (Sequential)	(None, 256)	66816
dense_8 (Dense)	multiple	2570

=====
Total params: 21,207,882
Trainable params: 21,206,858
Non-trainable params: 1,024
=====

In [20]:

```
import datetime as dt

earlystopping = tf.keras.callbacks.EarlyStopping(monitor = 'val_loss', patience =
3)

checkpoint = tf.keras.callbacks.ModelCheckpoint(filepath =
f'./saving/vgg_{dt.datetime.now().ctime()}.ckp',
                                                save_weights_only = True,
                                                save_best_only = True,
                                                monitor = 'val_loss',
                                                verbose = 1)
```

In [21]:

```
vgg_model2.compile(optimizer = 'adam',
                  loss = 'categorical_crossentropy',
                  metrics = ['acc'])
```

In [*]:

```
vgg_his2 = vgg_model2.fit(tr_gen,  
                          validation_data = val_gen,  
                          epochs = 30,  
                          callbacks = [checkpoint])
```

Epoch 1/30

240/469 [=====>.....] - ETA: 1:33:18 - loss: 2.9124 - acc: 0.
1409

In []:

```
plot_loss_acc(vgg_his2)
```

In []:

In []:

In []:

In []:

In []:

In []: